



```
static unsigned ifs_find_dirty_range(struct folio *folio,
                                   struct iomap_folio_state *ifs, u64 *range_start, u64 range_end)
{
    struct inode *inode = folio->mapping->host;
    unsigned start_blk =
        offset_in_folio(folio, *range_start) >> inode->i_blkbits;
    unsigned end_blk = min_not_zero(
        offset_in_folio(folio, range_end) >> inode->i_blkbits,
        i_blocks_per_folio(inode, folio));
    unsigned nblks = 1;

    while (!ifs_block_is_dirty(folio, ifs, start_blk))
        if (++start_blk == end_blk)
            return 0;

    while (start_blk + nblks < end_blk) {
        if (!ifs_block_is_dirty(folio, ifs, start_blk + nblks))
            break;
        nblks++;
    }

    *range_start = folio_pos(folio) + (start_blk << inode->i_blkbits);
    return nblks << inode->i_blkbits;
}

static unsigned iomap_find_dirty_range(struct folio *folio, u64 *range_start,
                                       u64 range_end)
{
    struct iomap_folio_state *ifs = folio->private;

    if (*range_start >= range_end)
        return 0;

    if (ifs)
        return ifs_find_dirty_range(folio, ifs, range_start, range_end);
    return range_end - *range_start;
}

unsigned f2fs_iomap_find_dirty_range(struct folio *folio, u64 *range_start,
                                     u64 range_end)
{

```

```
struct inode* inode=folio->mapping->host;
if(f2fs_compressed_file(inode))
{
    range_end=min(range_end,(range_end>>F2FS_I(inode)->i_log_cluster_size+1)<< F2FS_I(
}
return iomap_find_dirty_range(folio, range_start, range_end);
}
```

```

static int f2fs_write_single_folio(struct folio *folio, int *submitted_pages_count,
                                   struct bio **bio_ptr, sector_t *last_block_ptr,
                                   struct writeback_control *wbc,
                                   enum iostat_type io_type,
                                   unsigned int compr_blocks, bool is_reclaim,
                                   u64 dirty_pos_start, u64 dirty_pos_end)
{
    struct inode *inode = folio->mapping->host;
    struct f2fs_sb_info *sbi = F2FS_I_SB(inode);
    struct f2fs_iomap_folio_state *ifs = folio->private; // ifs should be set by caller if folio
    pgoff_t folio_start_index = folio->index;
    int err = 0;
    int local_submitted_count = 0;

    // Calculate the page indices within the folio that correspond to the dirty range
    // The dirty_pos_start/end are file offsets. We need folio-relative page indices.
    pgoff_t first_page_idx_in_folio = (dirty_pos_start - folio_pos(folio)) >> PAGE_SHIFT;
    pgoff_t last_page_idx_in_folio = (dirty_pos_end - 1 - folio_pos(folio)) >> PAGE_SHIFT;

    for (pgoff_t i = first_page_idx_in_folio; i <= last_page_idx_in_folio; ++i) {
        struct page *page = folio_page(folio, i);
        pgoff_t current_page_file_index = folio_start_index + i;

        // Skip if page is beyond EOF (i_size_read is expensive here,
        // writeback_iter and iomap_writepage_handle_eof should have handled major EOF iss
        // However, a check against wbc->range_end or a precise end_index might be useful
        // if not fully covered by iomap_find_dirty_range logic)
        if (current_page_file_index >= wbc_folio_end_index(wbc, folio)) {
            // This case should ideally be prevented by iomap_find_dirty_range
            // or iomap_writepage_handle_eof. If it occurs, it means the
            // dirty range extends beyond what wbc expects.
            continue;
        }

        // For large folios, increment pending bytes before attempting to write the page
        if (ifs) {
            atomic_add(PAGE_SIZE, &ifs->write_bytes_pending);
        }
    }
}

```

```

struct f2fs_io_info fio = {
    .sbi = sbi,
    .ino = inode->i_ino,
    .type = DATA, // Assuming DATA type, adjust if node pages are handled here
    .op = REQ_OP_WRITE,
    .op_flags = wbc_to_write_flags(wbc),
    .old_blkaddr = NULL_ADDR,
    .page = page,
    .encrypted_page = NULL,
    .submitted = 0, // f2fs_do_write_data_page will update this if it submits
    .compr_blocks = compr_blocks, // Pass this through
    .need_lock = compr_blocks ? LOCK_DONE : LOCK_RETRY,
    .meta_gc = f2fs_meta_inode_gc_required(inode) ? 1 : 0,
    .io_type = io_type,
    .io_wbc = wbc,
    // For bio and last_block, F2FS has its own merging.
    // If f2fs_do_write_data_page uses them directly, pass pointers.
    // If it's about merging *across calls* to f2fs_do_write_data_page,
    // that logic is more complex and currently sits higher up or in IPU.
    // For simplicity here, let's assume f2fs_do_write_data_page
    // handles its bio needs internally or via fio->bio.
    // The bio_ptr and last_block_ptr are from the original f2fs_write_single
    // which might have tried to merge. Here, each page is distinct for fio.
    .bio = bio_ptr ? *bio_ptr : NULL, // Allow passing bio for potential inter
    .last_block = last_block_ptr ? *last_block_ptr : 0,
};

// Retain original retry logic for f2fs_do_write_data_page
err = f2fs_do_write_data_page(&fio);
if (err == -EAGAIN) {
    f2fs_bug_on(sbi, compr_blocks && err == -EAGAIN); // Original bug_on
    fio.need_lock = LOCK_REQ;
    err = f2fs_do_write_data_page(&fio);
}

if (bio_ptr) // Update the caller's bio if f2fs_do_write_data_page modified it
    *bio_ptr = fio.bio;
if (last_block_ptr)
    *last_block_ptr = fio.last_block;

```

```

if (err) {
    // If this page failed to be prepared for write (e.g. -ENOSPC),
    // decrement pending_bytes if we incremented it.
    if (ifs) {
        // If submission failed for this page, it won't go to f2fs_write_e
        // for this portion.
        if (atomic_sub_and_test(PAGE_SIZE, &ifs->write_bytes_pending)) {
            // This implies all other pending IO for this folio also c
            // and this was the last one.
            // This scenario is tricky: if other pages *were* submitte
            // they will call folio_end_writeback. If this is the *only
            // page and it fails, then folio_end_writeback needs to be
            // by the caller (f2fs_write_cache_folios) after unlock.
            // The iomap model handles this by a bias count, which is
            // decremented after the loop. If it hits 0, and no IO was
            // then it calls folio_end_writeback.
            // For now, just decrement. The caller will handle overall
        }
    }
    // An error on one page within the folio typically means we stop
    // processing this folio and report the error.
    // The VFS writeback loop (writeback_iter) might then requeue the folio.
    mapping_set_error(folio->mapping, err); // Report error to mapping
    break; // Stop processing this folio on error
} else {
    // Page successfully processed by f2fs_do_write_data_page
    // fio.submitted should reflect if f2fs_do_write_data_page actually
    // queued something (it's usually 1 for a single page, or more if it did i
    // The original f2fs_write_single_data_page updated *submitted_count.
    // Here, we count pages for which f2fs_do_write_data_page succeeded.
    local_submitted_count++;
    if (wbc->nr_to_write > 0) // Only decrement if wbc cares
        wbc->nr_to_write--; // Decrement for each page successfully handle
}

// If wbc->nr_to_write is a budget and we've exhausted it.
if (wbc->nr_to_write == 0 && wbc->sync_mode == WB_SYNC_NONE) {

```

```

        err = -EAGAIN; // Signal to stop, budget met
        break;
    }
}

if (submitted_pages_count)
    *submitted_pages_count = local_submitted_count;

return err;
}
```C

```C
static int f2fs_write_cache_folios(struct address_space *mapping,
                                   struct writeback_control *wbc,
                                   enum iostat_type io_type)
{
    struct folio *folio =
        NULL;
    int err = 0;
    struct inode *inode = mapping->host;
    bool is_compressed_file = f2fs_compressed_file(inode);
    struct compress_ctx cc;
    int iter_err = 0; // Error from writeback_iter
#ifdef CONFIG_F2FS_FS_COMPRESSION
    if (is_compressed_file) {
        // Initialize compress_ctx structure fields (inode, sizes, etc.)
        cc.inode = inode;
        cc.log_cluster_size = F2FS_I(inode)->i_log_cluster_size;
        cc.cluster_size = 1 << cc.log_cluster_size;
        // cc.rpages and cc.cpages will be allocated by f2fs_init_compress_ctx
        cc.rpages = NULL;
        cc.cpages = NULL;
        cc.cluster_idx = NULL_CLUSTER;
        cc.nr_rpages = 0;
        cc.nr_cpages = 0;
        cc.valid_nr_cpages = 0;
        if (f2fs_init_compress_ctx(&cc)) {
            is_compressed_file = false;

```

```

    }
}
#endif

if (get_dirty_pages(mapping->host) <=
    SM_I(F2FS_M_SB(mapping))->min_hot_blocks)
    set_inode_flag(mapping->host, FI_HOT_DATA);
else
    clear_inode_flag(mapping->host, FI_HOT_DATA);

while ((folio = writeback_iter(mapping, wbc, folio, &iter_err))) {
    struct f2fs_iomap_folio_state *fifs = NULL;
    u64 pos = folio_pos(folio);
    u64 end_pos = pos + folio_size(folio);
    u64 end_pos_aligned;
    u32 r_len;
    int submitted_pages_this_folio = 0;
    if (!iomap_writepage_handle_eof(folio, inode, &end_pos)) {
        folio_unlock(folio);
        return 0;
    }

    fifs = folio->private;
    if (i_blocks_per_folio(inode, folio) > 1) {
        if (!fifs) {
            fifs = fifs_alloc(inode, folio, 0);
            iomap_set_range_dirty(folio, 0, end_pos - pos);
        }

        /*
         * Keep the I/O completion handler from clearing the writeback
         * bit until we have submitted all blocks by adding a bias to
         * fifs->write_bytes_pending, which is dropped after submitting
         * all blocks.
         */
        WARN_ON_ONCE(atomic_read(&fifs->write_bytes_pending) !=
            0);
        atomic_inc(&fifs->write_bytes_pending);
    }
    folio_start_writeback(folio);
}

```

```

end_aligned = round_up(end_pos, i_blocksize(inode));
while ((r_len = f2fs_iomap_find_dirty_range(folio, &pos,
                                           end_aligned))) {
    if (is_compressed_file)
    {
        pgoff_t cluster_idx =
            cluster_idx(&cc, pos >> PAGE_SHIFT);

        // If current dirty range belongs to a new cluster than what's in
        if (cc.cluster_idx != NULL_CLUSTER &&
            cc.cluster_idx != cluster_idx) {
            int submitted_compress = 0;
            if (!f2fs_cluster_is_empty(&cc)) {
                err = f2fs_write_multi_pages(
                    &cc,
                    &submitted_compress,
                    wbc, io_type);
                // wbc->nr_to_write -= submitted_compress;
                /* writeback_iter choose to dec folio_nr_pages in
                 * completely different from the dec policy that f2
                 * it's an open problem here
                 */
            }
            f2fs_destroy_compress_ctx(
                &cc,
                false); // false to reuse buffers, resets nr_rpages
            f2fs_init_compress_ctx(
                &cc); // Ensure rpages is valid
            cc.cluster_idx =
                NULL_CLUSTER; // Explicitly ensure for next logic
        }
        /*我们暂时没完全搞清prepare_compress_overwrite
        的防御性编程在f2fs_write_cache_folios中如何实现等价逻辑。先注释掉吧*/
        // // If cc is now for a new/empty cluster, prepare it
        // if (f2fs_cluster_is_empty(&cc)) {
        //     cc.cluster_idx =
        //         cluster_idx; // Set for prepare_compress_overwrite
        //     int prep_ret = prepare_compress_overwrite(

```



```

//          &cc, NULL,
//          pos >>
//          PAGE_SHIFT, /* pgoff_t index */
//          NULL);
//      if (prep_ret < 0) {
//          err = prep_ret;
//          goto handle_current_folio_error;
//      }
//      // After prepare_compress_overwrite, cc.cluster_idx is
//      // and cc.rpages might be filled if it was an overwrite
// }

// Add the current dirty range from folio to cc
// f2fs_compress_ctx_add_folio does not do folio_get.
// The folio is locked. cc needs a ref.
folio_get(folio); // cc takes a reference
f2fs_compress_ctx_add_folio(&cc, folio, pos, r_len);
// The ref taken by folio_get will be put by f2fs_put_rpages_wbc (
/*靠判断相邻两次cluster不连续 然后就进行f2fs_write_multi_pages的
那个逻辑,是足够覆盖掉cluster满的这个逻辑的*/
// if (f2fs_cluster_is_full(&cc)) {
//     int submitted_compress = 0;
//     err = f2fs_write_multi_pages(
//         &cc, &submitted_compress, wbc,
//         io_type);
//     // wbc->nr_to_write -= submitted_compress;
//     if (err) {
//         if (err == -EAGAIN)
//             wbc->nr_to_write = 0;
//         goto handle_current_folio_error;
//     }
//     f2fs_destroy_compress_ctx(
//         &cc,
//         false); // Reset for next potential cluster
//     f2fs_init_compress_ctx(&cc);
//     cc.cluster_idx = NULL_CLUSTER;
// }

}
else

```

```

        { // Non-compressed file path
            int submitted_this_range = 0;
            // For non-compressed, f2fs_write_single_folio handles one dirty range
            // It internally creates f2fs_io_info (fio) and calls f2fs_do_write_folio
            // It needs the folio (folio), and the dirty range.
            err = f2fs_write_single_folio(folio, &submitted_this_range, NULL, NULL);
            // submitted_pages_this_folio += submitted_this_range; // f2fs_write_folio

        }
        pos += r_len; // Advance for next iomap_find_dirty_range call
    } // End of while(iomap_find_dirty_range)

    // After processing all dirty ranges for this folio
    iomap_clear_range_dirty(folio, 0, folio_size(folio));

handle_current_folio_error: // Common error handling for the current folio
    folio_unlock(folio);
    if (fifs) {
        if (atomic_dec_and_test(&fifs->write_bytes_pending)) {
            folio_end_writeback(folio);
        }
    } else {
        if (!err &&
            pos ==
                current_folio_pos_in_file) { // No dirty ranges processed
            folio_end_writeback(folio);
        } else if (err && err != -EAGAIN) {
            // If an error stopped processing before any IO for this folio,
            // or if f2fs_write_single_folio failed before submitting.
            // This needs to be robust.
            // For now, rely on fifs or assume single-page folios' end_io handler
        }
    }
    if (err &&
        err != -EAGAIN) { // If a real error occurred for this folio
        mapping_set_error(mapping, err);
        // writeback_iter will stop if wbc->sync_mode != WB_SYNC_NONE on error.
        // Or we can break here.
        break;
    }
}

```

```

} // End of while(writeback_iter)

if (is_compressed_file && !f2fs_cluster_is_empty(&cc)) {
    int submitted_compress = 0;
    int flush_err = f2fs_write_multi_pages(&cc, &submitted_compress,
                                           wbc, io_type);

    // wbc->nr_to_write -= submitted_compress;
    if (flush_err && !err) { // Prioritize earlier errors
        err = flush_err;
    }
}

if (is_compressed_file) {
    // Free cc.rpages and cc.cpages if they were allocated
    f2fs_destroy_compress_ctx(
        &cc,
        true); // true to free cpages, rpages also freed by page_array_free
}

// If err was -EAGAIN from writeback_iter (iter_err) or our loop, VFS expects 0.
if (err == -EAGAIN || iter_err == -EAGAIN) {
    return 0;
}
return err ? err : iter_err;
}

```

```

static int prepare_compress_overwrite(struct compress_ctx *cc,
                                     struct page **pagep, pgoff_t index, void **fsdata)
{
    struct f2fs_sb_info *sbi = F2FS_I_SB(cc->inode);
    struct address_space *mapping = cc->inode->i_mapping;
    struct folio *folio;
    sector_t last_block_in_bio;
    fgf_t fgf_flag = FGP_LOCK | FGP_WRITE | FGP_CREAT;
    pgoff_t start_idx = start_idx_of_cluster(cc);
    int i, ret;

retry:
    ret = f2fs_is_compressed_cluster(cc->inode, start_idx);
    if (ret <= 0)
        return ret;

    ret = f2fs_init_compress_ctx(cc);
    if (ret)
        return ret;

    /* keep folio reference to avoid page reclaim */
    for (i = 0; i < cc->cluster_size; ) {
        folio = f2fs_filemap_get_folio(mapping, start_idx + i,
                                       fgf_flag, GFP_NOFS);
        if (IS_ERR(folio)) {
            ret = PTR_ERR(folio);
            goto unlock_pages;
        }
        bool needs_read = false;
        pgoff_t idx_in_folio=offset_in_folio(folio,(start_idx+i)<<PAGE_SHIFT)>>PAGE_SHIFT;
        unsigned int num_pages_iter=min(folio_nr_pages(folio)-idx_in_folio,cc->cluster_size-idx_in_folio);
        for (; idx_in_folio < num_pages_iter; ++idx_in_folio) {

            struct f2fs_iomap_folio_state* fifs=folio->private;
            if (!f2fs_ifs_block_is_uptodate(fifs, idx_in_folio)) {
                needs_read = true;
                break;
            }
        }
    }
}

```

```

        if(needs_read)
        {
            f2fs_compress_ctx_add_folio(cc,folio,(folio->index + idx_in_folio) << PAGE_SHIFT);
        }
        else
        {
            folio_put(folio);
        }
        i+=num_pages_iter;
    }

    if (!f2fs_cluster_is_empty(cc)) {
        struct bio *bio = NULL;

        ret = f2fs_read_multi_folios(cc, &bio, cc->cluster_size,
                                     &last_block_in_bio, NULL, true);
        f2fs_put_rpages(cc);
        f2fs_destroy_compress_ctx(cc, true);
        if (ret)
            goto out;
        if (bio)
            f2fs_submit_read_bio(sbi, bio, DATA);

        ret = f2fs_init_compress_ctx(cc);
        if (ret)
            goto out;
    }

    for (i = 0; i < cc->cluster_size; ) {
        f2fs_bug_on(sbi, cc->rpages[i]);

        folio = filemap_lock_folio(mapping, start_idx + i);
        if (IS_ERR(folio)) {
            /* folio could be truncated */
            goto release_and_retry;
        }

        f2fs_folio_wait_writeback(folio, DATA, true, true);
        pgoff_t idx_in_folio=offset_in_folio(folio,(start_idx+i)<<PAGE_SHIFT)>>PAGE_SHIFT;
    }

```

```

        unsigned int num_pages_to_add=min(folio_nr_pages(folio)-idx_in_folio,cc->cluster_s
        f2fs_compress_ctx_add_folio(cc, folio,(start_idx + i) >> PAGE_SHIFT,num_pages_to_a

        if (!folio_test_uptodate(folio)) {
            f2fs_handle_page_eio(sbi, folio, DATA);
release_and_retry:
            f2fs_put_rpages(cc);
            f2fs_unlock_rpages(cc, i + 1);
            f2fs_destroy_compress_ctx(cc, true);
            goto retry;
        }
        i+=num_pages_to_add;
    }

    if (likely(!ret)) {
        *fsdata = cc->rpages;
        *pagep = cc->rpages[offset_in_cluster(cc, index)];
        return cc->cluster_size;
    }

unlock_pages:
    f2fs_put_rpages(cc);
    f2fs_unlock_rpages(cc, i);
    f2fs_destroy_compress_ctx(cc, true);
out:
    return ret;
}

```

```

static void f2fs_drop_rpages(struct compress_ctx *cc, int len, bool unlock)
{
    int i;
    unsigned int num_to_skip=0;
    for (i = 0; i < len; ) {
        if (!cc->rpages[i])
        {
            num_to_skip = 1;
            continue;
        }
        struct folio* rfolio=page_folio(cc->rpages[i]);
        if (unlock)
            folio_unlock(rfolio);
        else
            folio_put(rfolio);
        num_to_skip =min(folio_nr_pages(rfolio)-folio_page_idx(rfolio,cc->rpages[i]), len-i);
        i+=num_to_skip;
    }
}

static void f2fs_put_rpages(struct compress_ctx *cc)
{
    f2fs_drop_rpages(cc, cc->cluster_size, false);
}

static void f2fs_unlock_rpages(struct compress_ctx *cc, int len)
{
    f2fs_drop_rpages(cc, len, true);
}

static void f2fs_put_rpages_wbc(struct compress_ctx *cc,
                                struct writeback_control *wbc, bool redirty, int unlock)
{
    int i;
    unsigned int num_to_skip=0;
    for (i = 0; i < len; ) {
        if (!cc->rpages[i])
        {
            num_to_skip = 1;

```

```
        continue;
    }
    struct folio* rfolio=page_folio(cc->rpages[i]);
    if (redirty)
        folio_redirty_for_write_page(rfolio);
    f2fs_folio_put(page_folio(cc->rpages[i]), unlock);
    num_to_skip =min(folio_nr_pages(rfolio)-folio_page_idx(rfolio,cc->rpages[i]), len-
    i+=num_to_skip;
}
}
```